

The head start your agent gets the day SolveWorks builds it.

Most businesses do not need another chatbot. They need an agent that understands the work: the sources, rules, approvals, tools, and corrections that make execution safe.

WHAT SETS SOLVEWORKS APART

We build the Company Brain around the agent.

Context layer. Workflow rules. Permission model. Skillset Packs. QA checks. Execution surfaces.

THE HARD LESSON

The model is rarely the whole problem.

Every real agent build since December has confirmed one thing: a better prompt does not fix a broken operating layer.

Cold starts

The agent wakes up without the right call notes, SOPs, dashboard definitions, CRM history, or priorities.

No source of truth

It cannot know whether the spreadsheet, dashboard, transcript, old SOP, or correction should win.

Unsafe permissions

Nobody has defined what the agent can read, draft, send, change, or escalate.

Lost corrections

Humans correct the same mistake again because feedback never becomes a reusable rule.

SolveWorks fixes the operating layer first, then gives the agent room to execute.

NOT A CHATBOT INSTALL

We build agents like operators, not demos.

A useful business agent needs more than model access. It needs trusted context, scoped tools, approval gates, feedback loops, logs, QA, and a way to improve after every correction.

Company Brain

Map the business context the agent needs before it acts.

Skillset Packs

Package how the company does recurring work.

Execution Surfaces

Turn work into briefs, dashboards, drafts, watchdogs, and queues.

Safety + Receipts

Start safe, log what happened, and expand only after verification.

THE FOUNDATION

The agent starts smarter because the brain is designed.

We do not dump the whole company into a prompt. We design the layers that let an agent retrieve the right context and respect the business rules.

The Company Brain is the operating layer between scattered business knowledge and useful agent execution.

It defines what the agent should know, which source to trust, what it is allowed to do, and how corrections improve the next run.

This is the difference between an agent that impresses in a demo and one that earns a place in daily operations.

SIX LAYERS, PART ONE

First: capture the work and retrieve the right context.

The agent should not start cold, and it should not be flooded with everything. It needs the right facts for the job in front of it.

01

Capture

Calls, SOPs, docs, dashboards, CRM notes, decisions, corrections.

02

Retrieval

Only the facts needed for the job, not a noisy memory dump.

03

Source Truth

Rules for which system wins when data conflicts.

SIX LAYERS, PART TWO

Then: set boundaries, learn from corrections, and execute.

This is where the agent becomes safer and more useful. Permissions are explicit, feedback becomes reusable, and outputs move into the tools the business actually uses.

04

Permissions

What the agent can read, draft, send, change, or escalate.

05

Feedback Loops

Corrections become future rules, not lost chat threads.

06

Execution

Briefs, dashboards, follow-ups, watchdogs, tasks, and approvals.

HOW AGENTS LEARN THE WORK

We package the way your company actually does the job.

The strongest agents do not rely on vague instructions. They run on reusable operating knowledge: decisions, checks, commands, and examples.

INCLUDED

Rules

How to decide, what to avoid, when to escalate, and which source is trusted.

INCLUDED

Runbooks

The step-by-step operating pattern for repeatable workflows.

INCLUDED

QA checks

Tests, review gates, leakage scans, and receipts before work is trusted.

INCLUDED

Tool commands

Safe ways the agent reads, drafts, checks, reports, and operates systems.

Our builds are designed to compound: each workflow makes the next one faster to stand up and easier to improve.

THE FIRST USEFUL WIN

Start where the team keeps re-explaining the work.

The goal is not “AI everywhere.” The goal is one workflow where context, judgment, and follow-through keep slowing the business down every week.

The best first workflow has three traits.

It repeats often. It depends on context. It benefits from a human approval step before action.

The starting point is practical: read-only or draft-only first, then approval-gated action once the workflow is proven.

COMMON FIRST PILOTS

Useful agents usually start with ordinary recurring work.

Sales follow-up

Call summary, buyer pain, next steps, CRM note, approved follow-up draft.

Owner briefing

Calendar, inbox, dashboard exceptions, priorities, risks, and decisions.

Dashboard analyst

Freshness checks, KPI explanations, anomalies, and recommended actions.

Ops watchdog

Recurring checks, stale tasks, order/inventory exceptions, and escalations.

BEFORE PROPOSING

A strong agent is scoped before it is built.

The call is not a generic demo. We map the workflow, sources, decisions, permissions, and risk boundaries so the proposal is specific.

1

What work repeats every week?

Find the workflow with the fastest practical return.

2

Which systems does it depend on?

Identify the calls, docs, dashboards, CRM, inboxes, and tools.

3

Which source should win?

Define source truth before the agent trusts bad data.

BEFORE PROPOSING

We also define what the agent should not do.

Good scoping protects the business. We decide what stays human-approved, what repeated mistakes become rules, and what first output would be useful immediately.

4

What must stay human-approved?

Set the permission ladder and escalation rules.

5

What mistakes repeat?

Turn recurring corrections into rules and QA checks.

6

What output would be useful tomorrow?

Pick the first brief, draft, dashboard, task list, or watchdog.

WHAT A SOLVEWORKS BUILD GIVES THE AGENT

Your agent does not start from a blank chat box.

The first build is not just access to a model. It is a working operating package around the agent so it knows the job, tools, boundaries, and definition of done.

Brain map

Sources, decisions, and context the workflow depends on.

Skillset Pack

Rules, prompts, runbooks, examples, checks, and escalation paths.

Tool bench

Safe ways to read, draft, report, summarize, check, and queue work.

QA receipts

Tests, launch checks, logs, and proof before permissions expand.

This is the head start: the agent begins with operating knowledge, not guesswork.

BRIEF BIO

SolveWorks builds practical AI agents for real business workflows.

We help owner-led companies turn scattered knowledge, repeated decisions, and manual follow-through into agent-assisted operating systems.

What we do

Build purpose-built agents, dashboards, briefs, watchdogs, and workflow tools.

How we work

Start with one useful workflow, map the sources and approvals, then launch safely.

What makes it different

The agent gets a Company Brain, Skillset Pack, tool bench, permission ladder, and feedback loop.

Who it is for

Business owners and operators who want AI help with follow-up, reporting, research, dashboards, admin, and recurring operations.

WHAT HAPPENS NEXT

When there is a fit, we send a specific proposal.

The proposal is built around the workflow we mapped together. It defines what the agent reads, what it produces, what needs approval, and how we verify that it works.

Operating model

Sources, source-truth rules, permission ladder, escalation path.

Working artifacts

Agent workflows, briefs or dashboards, runbooks, QA checks, receipts.

Launch posture

Safe first: read-only or draft-only before approved action.

Compounding value

Each correction improves the brain instead of disappearing in chat.

A useful agent is the model plus the business context, trust rules, permissions, skillsets, tools, QA, and feedback loops around it.